Endpoint Scanner: trace any ViT response back to the input

One scalar target in the real model → one input-coordinate learning image

Real input · sloth bear



1 · Choose a scalar response

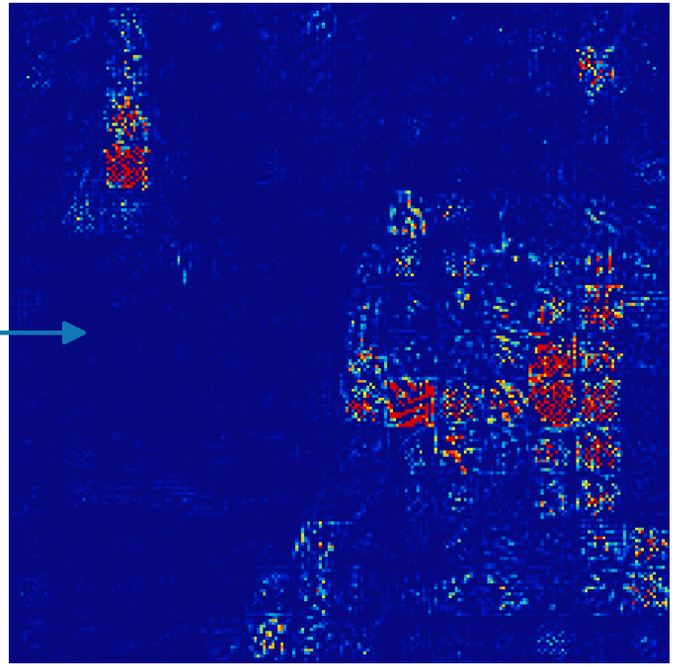
Ground-truth class logit
y = 297 · n02134418

2 · Differentiate at the endpoint

$$E_s(i) = x_i \, \partial s(x)/\partial x_i$$

The same input coordinates are retained
for logits, heads, branches, and internal targets.

Ground-truth class logit



input

Endpoint rule

where the positive response comes from

Main view: max(E, 0), per-panel P99 · full signed map retained in the supplement