Three related methods, three different questions

Architecture supplies the Jacobian; the seed and evaluation rule define the attribution

Original NNS-CNN

$$x \odot \nabla s(x)$$

Current ReLU gates and MaxPool winner

CNN endpoint rule

ViT Endpoint Scanner

$$x \odot \nabla s(x)$$

Full ViT Jacobian: LN, Q/K/V, residual, MLP

Our full-chain scan

Path-integrated attribution

$$\Delta x \odot \int_0^1 \nabla s(x_\alpha) d\alpha$$

Average gradients along a specified baseline path

Different evaluation rule

Endpoint identity does not imply completeness; Path IG additionally requires a baseline and a path.