Endpoint Scanner: from registered targets to behavioral evidence

A two-model, 400-image endpoint study with separate identity, correctness, and utility layers

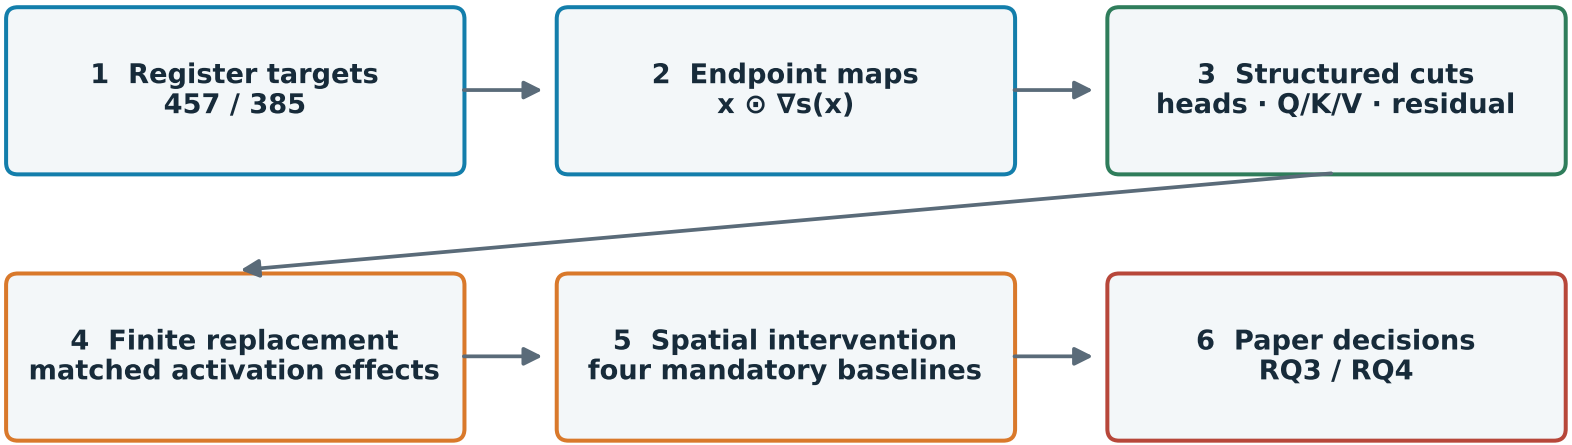

Identity Endpoint = Gradient × Input for the same scalar target

Correctness 76,800 closure checks · 28,800 negative controls · 6,144 analytic JVP checks

Utility RQ3 overall supported (11/12) · RQ4 not supported

Path IG: ViT-B calibration failed; no cross-model Path claim